

9/10/25

10. MNIST compression using Autoencoders

Aim:

To perform image compression on MNIST dataset
using an Autoencoder

Algorithm:

- * Load and Preprocess data:

Load MNIST dataset of handwritten digits

Normalize pixel values to range $[0,1]$

- * Build the Autoencoder architecture

Define the Encoder: reduce input using Dense layers

Define the Decoder: reconstruct the input image

from compressed latent vector

- * Compile the Autoencoder:

Use Adam Optimizer and binary cross entropy loss

- * Train on MNIST training data using $\text{input} = \text{output}$

- * Encode and Decode

Encoder to generate compressed representation

Decoder to reconstruct images from these compressed representations.

Pseudocode:

BEGIN

LOAD MNIST dataset

NORMALIZE all pixel values to $[0,1]$

FLATTEN each image into 784-dimensional vector

SET encoding-dim = 32

INPUT-LAYER = Input(shape = (784))

ENCODE1 = Dense(128, activation = "relu") (INPUT-LAYER)

ENCODE2 = Dense(64, activation = "relu") (ENCODE1)

ENCODE3 = Dense(encoding-dim, activation = "relu") (ENCODE2)

DECODE1 = Dense(64, activation = "relu") (ENCODE3)

DECODE2 = Dense(128, activation = "relu") (DECODE1)

CREATE autoencoder MODEL (INPUT-LAYER \rightarrow OUTPUT LAYER)

COMPILE MODEL using Adam optimizer and
binary-crossentropy loss

TRAIN autoencoder on (x_train, x_train)

for 50 epochs with batch-size = 256

ENCODE test images using encoder Part

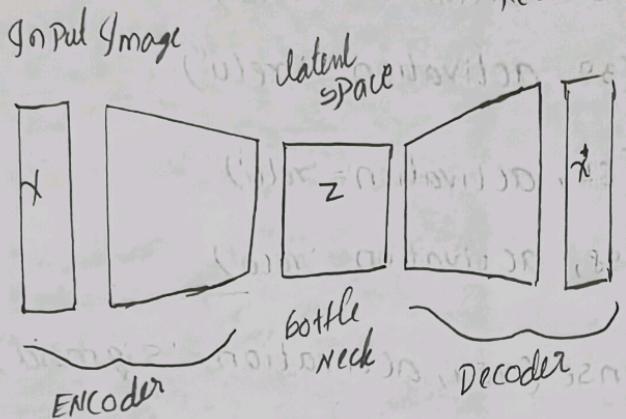
DECODE them back using autoencoder

DISPLAY original and reconstructed images

~~COMPUTE compression-ratio = 784/encoding-dim~~

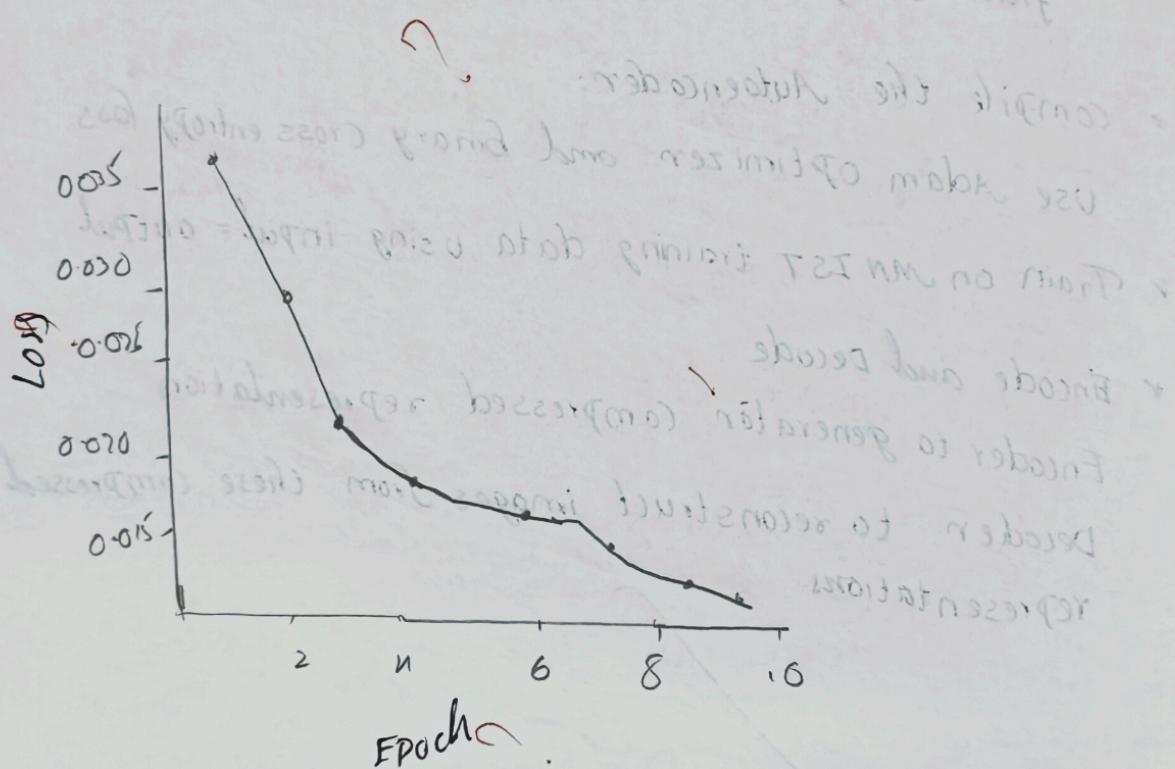
END

Autoencoder Architecture



EPOCH(1) : loss : 0.0378
 EPOCH(2) : loss : 0.0290
 EPOCH(3) : loss : 0.0229
 EPOCH(4) : loss : 0.149
 EPOCH(5) : loss : 0.143
 Epoch (8) : loss : 0.0137
 Epoch (9) : loss : 0.0133
 Epoch (10) : loss : 0.0133

Training loss over epochs



Result:
 Implementation of compression using Autoencoders was
 done successfully

11. Variational Auto Encoders

Aim:

To implement and analyze the Performance of a variational Autoencoder for learning compressed latent representations of data.

Algorithm:

- * Input: Training dataset x
- * Initialize: Encoder and Decoder neural networks

* For each epochs

Pass input x through Encoder \rightarrow get μ, σ

sample $z = \mu + \sigma * \epsilon$

Pass z through Decoder \rightarrow get reconstruction x'

compute:

$$\text{Reconstruction loss} = \|x - x'\|^2$$

$$\text{KL divergence} = -0.5 \sum (1 + \log(\sigma^2) - \mu^2 - \sigma^2)$$

$$\text{Total loss} = \text{Reconstruction loss} + \text{KL divergence}$$

After training, Generate new samples

Pseudocode:

Input: Dataset x

set learning_rate = 0.001

set latent_dimension = 6

Initialize Encoder network parameters (w_{enc}, b_{enc})

Initialize Decoder network parameters (w_{dec}, b_{dec})

FUNCTION ENCODER(x):

$$h = \text{ReLU}(w_1 * x + b_1)$$

$$\mu = w_\mu * h + b_\mu$$

$$\log \sigma^2 = w_\sigma * h + b_\sigma$$

return $\mu, \log \sigma^2$

FUNCTION DECODER(z)

$$h = \text{ReLU}(w_3 * z + b_3)$$

return x

For each epoch in range(num_epochs):

For each batch (x_batch) in Dataset x :

$$\mu = \log \sigma^2 = \text{ENCODER}(x_batch)$$

$$z = \text{SAMPLE}(\mu, \log \sigma^2)$$

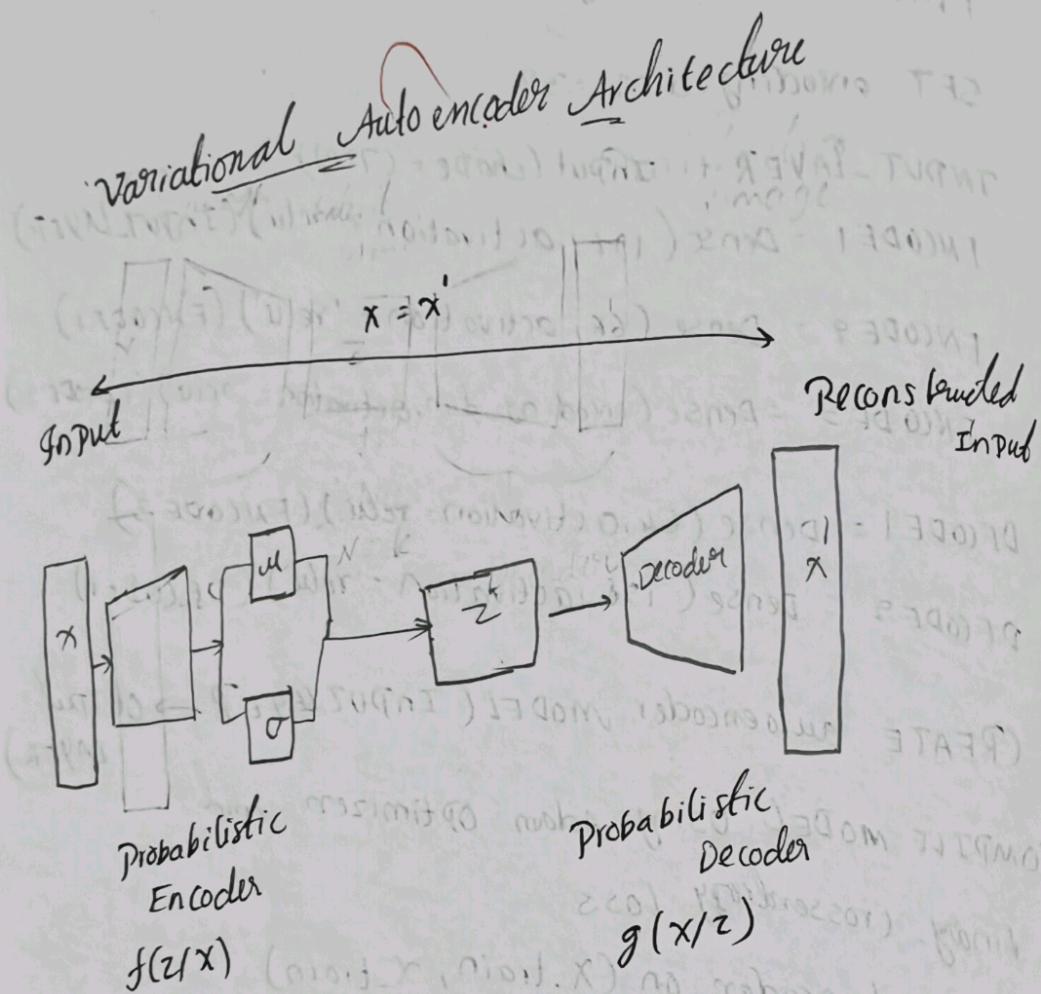
~~$$x' = \text{DECODER}(z)$$~~

~~$$\text{loss} = \text{loss}(x_batch, x', \mu, \log \sigma^2)$$~~

Backpropagate loss

Update encoder and decoder weights

FUNCTION Generate():



Output

EPOCH 1 : LOSS : 164.2524

EPOCH 2 : LOSS : 121.3644

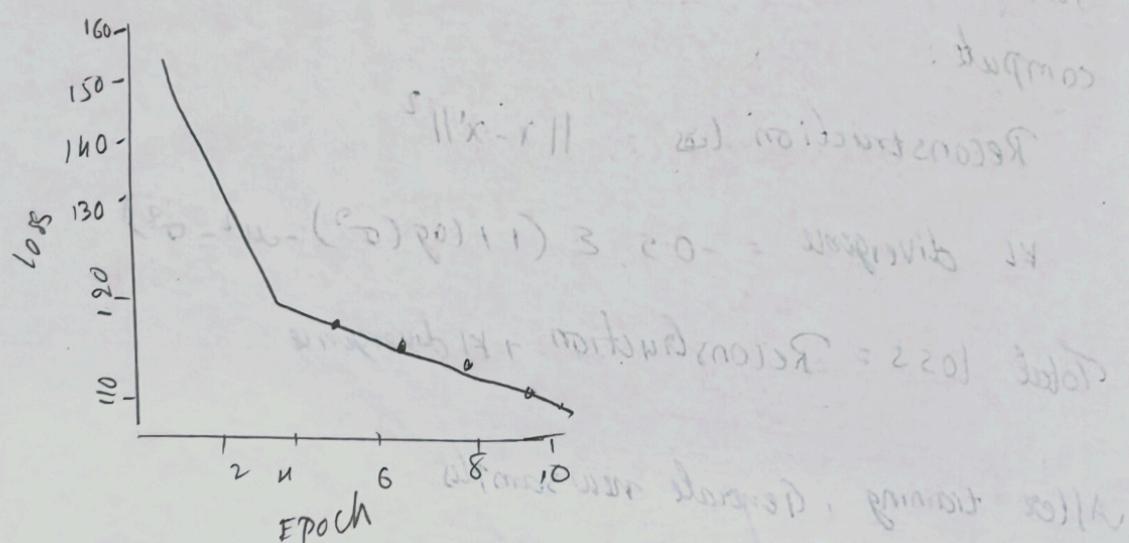
EPOCH 3 : LOSS : 114.3897

EPOCH 4 : LOSS : 111.3462

EPOCH 5 : LOSS : 109.5547

Avg LOSS : 111.1155

Training over loss over epochs



Result : Implementation of variation Autoencoders was done successfully

```

import torch
import torch.nn as nn
import torch.nn.functional as F
from torchvision import datasets, transforms
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt

# Encoder
class Encoder(nn.Module):
    def __init__(self, input_dim=784, hidden_dim=400, latent_dim=20):
        super(Encoder, self).__init__()
        self.fc1 = nn.Linear(input_dim, hidden_dim)
        self.fc_mu = nn.Linear(hidden_dim, latent_dim)
        self.fc_logvar = nn.Linear(hidden_dim, latent_dim)

    def forward(self, x):
        h = F.relu(self.fc1(x))
        mu = self.fc_mu(h)
        logvar = self.fc_logvar(h)
        return mu, logvar

# Decoder
class Decoder(nn.Module):
    def __init__(self, latent_dim=20, hidden_dim=400, output_dim=784):
        super(Decoder, self).__init__()
        self.fc1 = nn.Linear(latent_dim, hidden_dim)
        self.fc2 = nn.Linear(hidden_dim, output_dim)

    def forward(self, z):
        h = F.relu(self.fc1(z))
        x_recon = torch.sigmoid(self.fc2(h))
        return x_recon

# VAE
class VAE(nn.Module):
    def __init__(self):
        super(VAE, self).__init__()
        self.encoder = Encoder()
        self.decoder = Decoder()

    def reparameterize(self, mu, logvar):
        std = torch.exp(0.5 * logvar)
        eps = torch.randn_like(std)
        return mu + eps * std

    def forward(self, x):
        mu, logvar = self.encoder(x)
        z = self.reparameterize(mu, logvar)
        x_recon = self.decoder(z)
        return x_recon, mu, logvar

# Loss Function
def loss_function(x_recon, x, mu, logvar):
    BCE = F.binary_cross_entropy(x_recon, x, reduction='sum')
    KLD = -0.5 * torch.sum(1 + logvar - mu.pow(2) - logvar.exp())
    return BCE + KLD

# Data
transform = transforms.ToTensor()
train_data = datasets.MNIST(root='./data', train=True, transform=transform, download=True)
train_loader = DataLoader(train_data, batch_size=128, shuffle=True)

# Train
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
vae = VAE().to(device)
optimizer = torch.optim.Adam(vae.parameters(), lr=1e-3)

for epoch in range(5):
    total_loss = 0
    for x, _ in train_loader:
        x = x.view(-1, 784).to(device)
        x_recon, mu, logvar = vae(x)
        loss = loss_function(x_recon, x, mu, logvar)
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
        total_loss += loss.item()
    print(f"Epoch {epoch+1}, Loss: {total_loss/len(train_loader.dataset):.4f}")

# Generate new images
with torch.no_grad():

```

```

z = torch.randn(16, 20).to(device)
samples = vae.decoder(z).cpu().view(-1, 1, 28, 28)

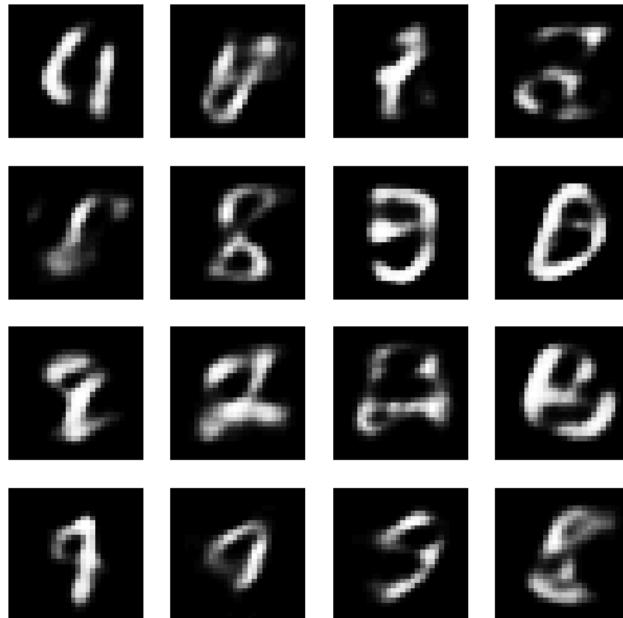
plt.figure(figsize=(6,6))
for i in range(16):
    plt.subplot(4,4,i+1)
    plt.imshow(samples[i].squeeze(), cmap='gray')
    plt.axis('off')
plt.show()

```

```

100%|██████████| 9.91M/9.91M [00:01<00:00, 6.09MB/s]
100%|██████████| 28.9k/28.9k [00:00<00:00, 160kB/s]
100%|██████████| 1.65M/1.65M [00:01<00:00, 1.52MB/s]
100%|██████████| 4.54k/4.54k [00:00<00:00, 12.0MB/s]
Epoch 1, Loss: 165.0683
Epoch 2, Loss: 122.3628
Epoch 3, Loss: 114.9047
Epoch 4, Loss: 111.8477
Epoch 5, Loss: 110.0657

```



```

import torch
import torch.nn as nn
import torch.nn.functional as F
from torchvision import datasets, transforms
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt

# Encoder
class Encoder(nn.Module):
    def __init__(self, input_dim=784, hidden_dim=400, latent_dim=20):
        super(Encoder, self).__init__()
        self.fc1 = nn.Linear(input_dim, hidden_dim)
        self.fc_mu = nn.Linear(hidden_dim, latent_dim)
        self.fc_logvar = nn.Linear(hidden_dim, latent_dim)

    def forward(self, x):
        h = F.relu(self.fc1(x))
        mu = self.fc_mu(h)
        logvar = self.fc_logvar(h)
        return mu, logvar

# Decoder
class Decoder(nn.Module):
    def __init__(self, latent_dim=20, hidden_dim=400, output_dim=784):
        super(Decoder, self).__init__()
        self.fc1 = nn.Linear(latent_dim, hidden_dim)
        self.fc2 = nn.Linear(hidden_dim, output_dim)

    def forward(self, z):
        h = F.relu(self.fc1(z))
        x_recon = torch.sigmoid(self.fc2(h))
        return x_recon

# VAE
class VAE(nn.Module):
    def __init__(self):

```

```

super(VAE, self).__init__()
self.encoder = Encoder()
self.decoder = Decoder()

def reparameterize(self, mu, logvar):
    std = torch.exp(0.5 * logvar)
    eps = torch.randn_like(std)
    return mu + eps * std

def forward(self, x):
    mu, logvar = self.encoder(x)
    z = self.reparameterize(mu, logvar)
    x_recon = self.decoder(z)
    return x_recon, mu, logvar

# Loss Function
def loss_function(x_recon, x, mu, logvar):
    BCE = F.binary_cross_entropy(x_recon, x, reduction='sum')
    KLD = -0.5 * torch.sum(1 + logvar - mu.pow(2) - logvar.exp())
    return BCE + KLD

# Data
transform = transforms.ToTensor()
train_data = datasets.MNIST(root='./data', train=True, transform=transform, download=True)
train_loader = DataLoader(train_data, batch_size=128, shuffle=True)

# Train
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
vae = VAE().to(device)
optimizer = torch.optim.Adam(vae.parameters(), lr=1e-3)

train_losses = []

for epoch in range(5):
    total_loss = 0
    for x, _ in train_loader:
        x = x.view(-1, 784).to(device)
        x_recon, mu, logvar = vae(x)
        loss = loss_function(x_recon, x, mu, logvar)
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
        total_loss += loss.item()
    avg_loss = total_loss / len(train_loader.dataset)
    train_losses.append(avg_loss)
    print(f"Epoch {epoch+1}, Loss: {avg_loss:.4f}")

# --- Plot training loss ---
plt.figure(figsize=(6,4))
plt.plot(train_losses, marker='o')
plt.title("VAE Training Loss")
plt.xlabel("Epoch")
plt.ylabel("Average Loss")
plt.grid(True)
plt.show()

# --- Reconstruction visualization ---
vae.eval()
with torch.no_grad():
    x, _ = next(iter(train_loader))
    x = x.view(-1, 784).to(device)
    x_recon, _, _ = vae(x)
    x = x.view(-1, 1, 28, 28)
    x_recon = x_recon.view(-1, 1, 28, 28)

# Show first 10 original and reconstructed images
n = 10
plt.figure(figsize=(15,4))
for i in range(n):
    # Original
    plt.subplot(2, n, i + 1)
    plt.imshow(x[i].cpu().squeeze(), cmap='gray')
    plt.title("Original")
    plt.axis('off')
    # Reconstructed
    plt.subplot(2, n, i + 1 + n)
    plt.imshow(x_recon[i].cpu().squeeze(), cmap='gray')
    plt.title("Reconstructed")
    plt.axis('off')
plt.show()

# --- Generate new samples ---

```

```
with torch.no_grad():
    z = torch.randn(16, 20).to(device)
    samples = vae.decoder(z).cpu().view(-1, 1, 28, 28)

plt.figure(figsize=(6,6))
for i in range(16):
    plt.subplot(4,4,i+1)
    plt.imshow(samples[i].squeeze(), cmap='gray')
    plt.axis('off')
plt.suptitle("Generated New Digits")
plt.show()
```

```
Epoch 1, Loss: 165.2717
Epoch 2, Loss: 122.1188
Epoch 3, Loss: 115.0454
Epoch 4, Loss: 112.0027
Epoch 5, Loss: 110.1681
```

