

Week-17

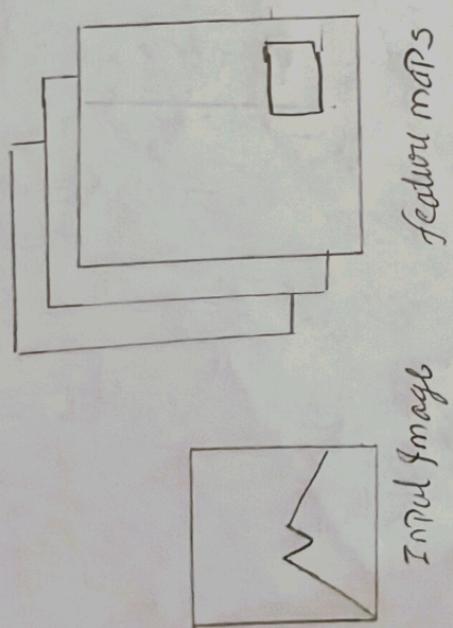
Aim: To implement a preTrained Convolutional Neural network such as ResNet50

Algorithm.

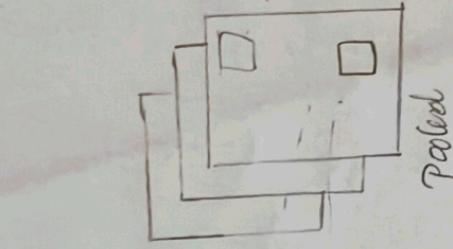
- * Import Libraries - PyTorch, Torchvision
- * Load a Pretrained CNN
- * Freeze all convolutional layers so that learned weights are not updated.
- * Replace the final classification layer with a new fully connected layer
- * Load and preprocess dataset
- * Train only the new classifier using features extracted by frozen CNN base
- * Evaluate model performance on test data

Observation

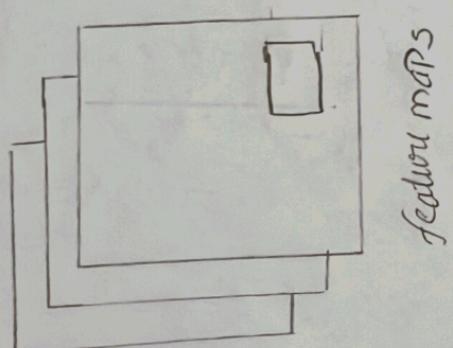
- * Faster Training, since the preTrained CNN layers are already trained on large dataset.
- * Improved Accuracy, Using a PreTrained model typically improves accuracy over a shallow model.
- * The lower layers of CNNs capture universal image features, which are useful across different image



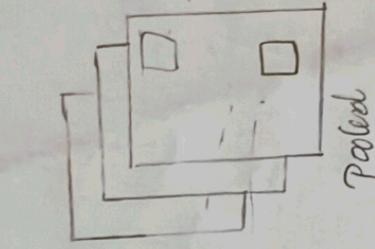
feature maps



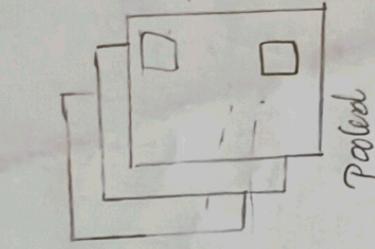
feature maps



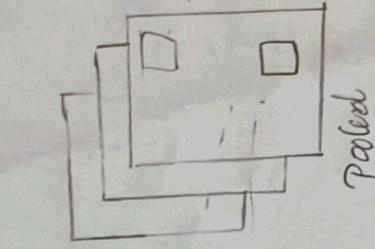
Pooled



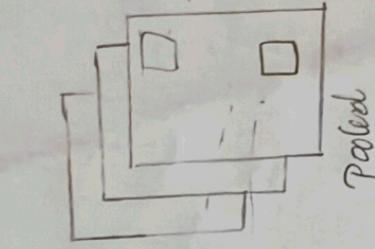
ReLU



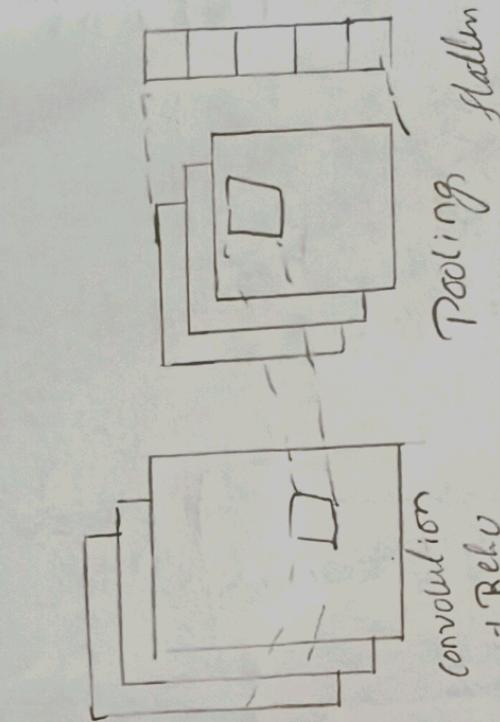
convolution



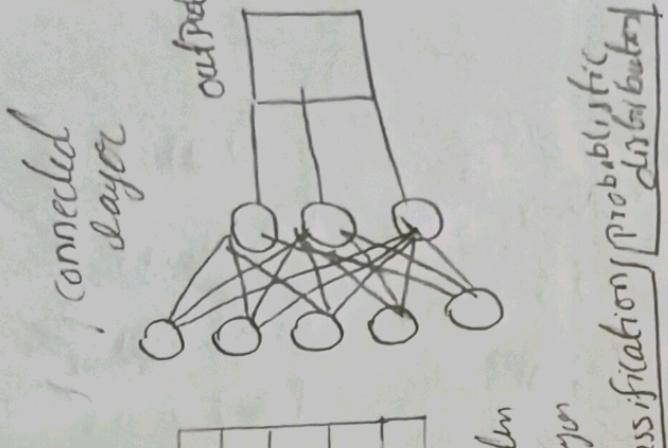
Pooling



softmax



output



classification / probability distribution

THAT Pretrained Architecture Model

Pseudocode

BEGIN

IMPORT torch, torchvision and required modules

LOAD PreTrained CNN model

for each Parameter in CNN base:

SET require_grad = False

REPLACE final layer with new linear layer

DEFINE loss function (crossEntropyLoss) and optimizer

LOAD Dataset and apply transformations (resize)

for each epoch:

for each batch (image, labels):

EXTRACT features using CNN base

PREDICT using classifier layer

compute loss and backpropagate

UPDATE weights of classifier layer only

Evaluate model accuracy on test set

END

Observation

- * Training time greatly reduced compared to training from scratch
- * The Model Achieves high Accuracy

Result

Pre trained model of CNN successfully implemented

~~15/01~~

Output

EPOCH [1/5] : loss : 0.848'

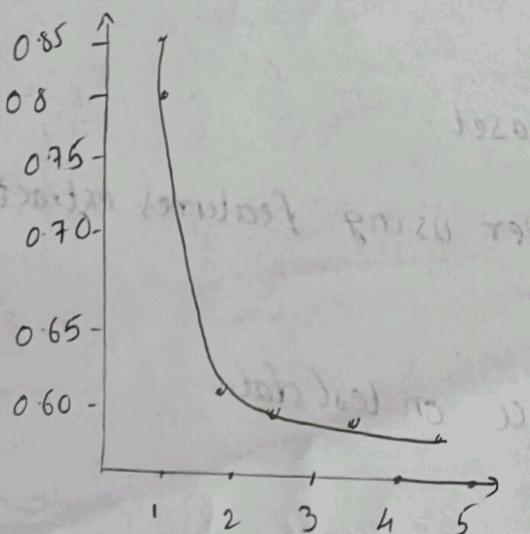
EPOCH [2/5] : loss : 0.6266

EPOCH [3/5] : loss : 0.5937

EPOCH [4/5] : loss : 0.5799

EPOCH [5/5] : loss : 0.5699

Training loss curve + Transfer learning
Parameter



```
# Experiment 14: Transfer Learning with MobileNetV2 (Fixed Version using CIFAR-10)
```

```
import tensorflow as tf

from tensorflow.keras.applications import MobileNetV2

from tensorflow.keras.layers import Dense, GlobalAveragePooling2D

from tensorflow.keras.models import Model


# Load CIFAR-10 dataset (10 classes)

(x_train, y_train), (x_test, y_test) = tf.keras.datasets.cifar10.load_data()

# Normalize pixel values

x_train, x_test = x_train / 255.0, x_test / 255.0


# Convert labels to one-hot encoding

y_train = tf.keras.utils.to_categorical(y_train, 10)

y_test = tf.keras.utils.to_categorical(y_test, 10)


# Load pre-trained MobileNetV2 (no top layer)

base_model = MobileNetV2(weights='imagenet', include_top=False,
input_shape=(32,32,3))

base_model.trainable = False # freeze base


# Add new classifier

x = base_model.output

x = GlobalAveragePooling2D()(x)

x = Dense(128, activation='relu')(x)

predictions = Dense(10, activation='softmax')(x)
```

```
model = Model(inputs=base_model.input, outputs=predictions)

# Compile model

model.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])

# Train the model

history = model.fit(x_train, y_train, validation_data=(x_test, y_test), epochs=5,
batch_size=64)

# Evaluate

loss, acc = model.evaluate(x_test, y_test)

print(f"\n Validation Accuracy: {acc:.2f}")
```

```

| # Experiment 14: Transfer Learning with MobileNetV2 (Fixed Version using CIFAR-10)

import tensorflow as tf
from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras.layers import Dense, GlobalAveragePooling2D
from tensorflow.keras.models import Model

# Load CIFAR-10 dataset (10 classes)
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.cifar10.load_data()

# Normalize pixel values
x_train, x_test = x_train / 255.0, x_test / 255.0

# Convert labels to one-hot encoding
y_train = tf.keras.utils.to_categorical(y_train, 10)
y_test = tf.keras.utils.to_categorical(y_test, 10)

# Load pre-trained MobileNetV2 (no top layer)
base_model = MobileNetV2(weights='imagenet', include_top=False, input_shape=(32,32,3))
base_model.trainable = False # freeze base

# Add new classifier
x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(128, activation='relu')(x)
predictions = Dense(10, activation='softmax')(x)

# Compile model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Train the model
history = model.fit(x_train, y_train, validation_data=(x_test, y_test), epochs=5, batch_size=64)

# Evaluate
loss, acc = model.evaluate(x_test, y_test)
print(f"\n Validation Accuracy: {acc:.2f}")

```

```

Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170498071/170498071 5s 0us/step
/tmp/ipython-input-208358661.py:19: UserWarning: `input_shape` is undefined or non-square, or `rows` is not in [96, 128, 160, 192, 224].
  base_model = MobileNetV2(weights='imagenet', include_top=False, input_shape=(32,32,3))
Epoch 1/5
782/782 41s 32ms/step - accuracy: 0.2567 - loss: 2.0496 - val_accuracy: 0.3212 - val_loss: 1.8752
Epoch 2/5
782/782 7s 9ms/step - accuracy: 0.3289 - loss: 1.8464 - val_accuracy: 0.3321 - val_loss: 1.8412
Epoch 3/5
782/782 7s 9ms/step - accuracy: 0.3401 - loss: 1.8144 - val_accuracy: 0.3412 - val_loss: 1.8223
Epoch 4/5
782/782 7s 9ms/step - accuracy: 0.3506 - loss: 1.7849 - val_accuracy: 0.3411 - val_loss: 1.8132
Epoch 5/5
782/782 7s 8ms/step - accuracy: 0.3614 - loss: 1.7643 - val_accuracy: 0.3510 - val_loss: 1.8017
313/313 12s 6ms/step - accuracy: 0.3510 - loss: 1.7954

Validation Accuracy: 0.35

```