

Week-11 RNN

Aim: To implement a Recurrent Neural Network using PyTorch that processes sequential data and makes predictions based on past info

Algorithm:

- * Import Libraries: Load PyTorch Modules
- * Prepare Dataset: Define a small sequence dataset & convert inputs and targets into tensor formats
- * Define RNN
 - Use nn.Embedding to map indices to dense vectors.
 - Add nn.RNN layer to capture sequential dependencies.
 - Add nn.Linear layer to map hidden state
- * Initialize Hyper Parameters:
 - vocabulary size, hidden layer size, learning rate, no of epochs.
 - Initialize optimizer and loss function
- * Train The Model, for each epoch
 - Initialize ~~hidden~~ state, Forward Pass: input \rightarrow RNN \rightarrow Output
 - complete loss b/w predicted output and target
 - Backpropagate gradients
- * Evaluate Model:
 - Feed the input sequence to trained model
 - Decode Predictions into readable formats

Pseudo code

BEGIN

Import torch, nn, optim

Step-1: Dataset Preparation

Define vocabulary and mapping (char2idx, idx2char)

Encode input sequence x,y

Convert x,y to tensor

Step-2: Define Model

CLASS RNNModel:

INIT (vocab-size, hidden-size):

embedding = nn.Embedding(vocab-size, hidden-size)

rnn = nn.RNN(hidden-size, hidden-size)

FORWARD (x, hidden)

x_embedded = embedding(x)

output, hidden = rnn(x_embedded, hidden)

Prediction = fc(output)

RETURN Prediction, hidden

Step-4 - Training loop

For epoch in range(num-epochs):

hidden = zeroes(initial-hidden-state)

output, hidden = model(x, hidden)

loss = loss_fn(output, y)

optimizer.zero_grad()

~~loss.backward()~~

Step-5 - Final evaluation

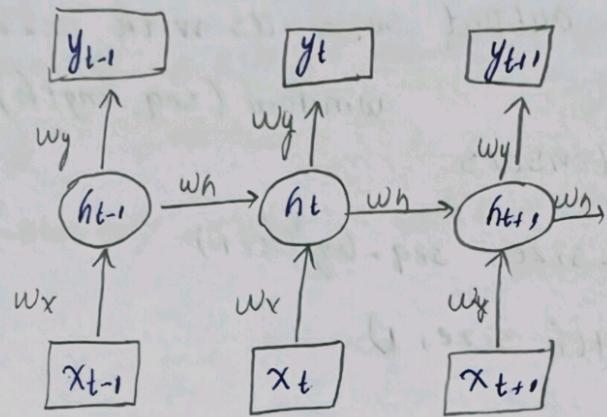
hidden = zeroes(initial-hidden-state)

output, hidden = model(x, hidden)

~~predicted-sequence = decode(output)~~

~~print(predicted-sequence)~~

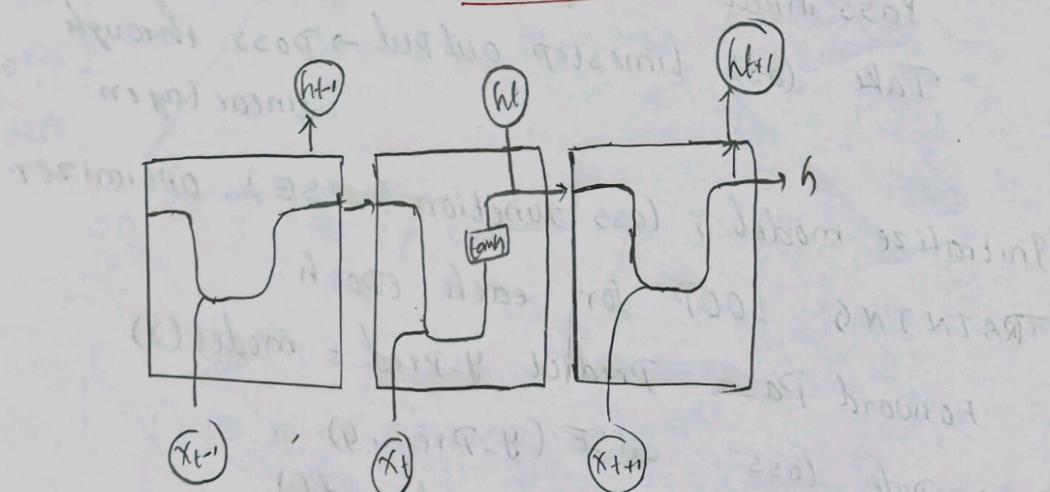
END



$$h_t = \tanh(w_h h_{t-1} + x_t w_x + b_h)$$

$$y_t = w_y h_t + b_y$$

RNN Internal Architecture



EPOCH 1
accuracy: 0.5433 - loss: 0.6817

EPOCH 2
accuracy: 0.7579 - loss: 0.5022

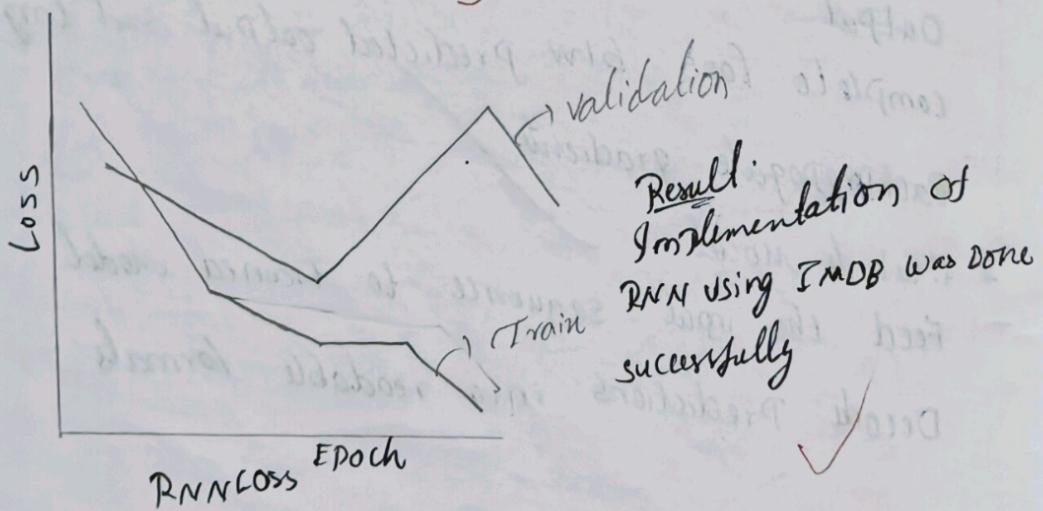
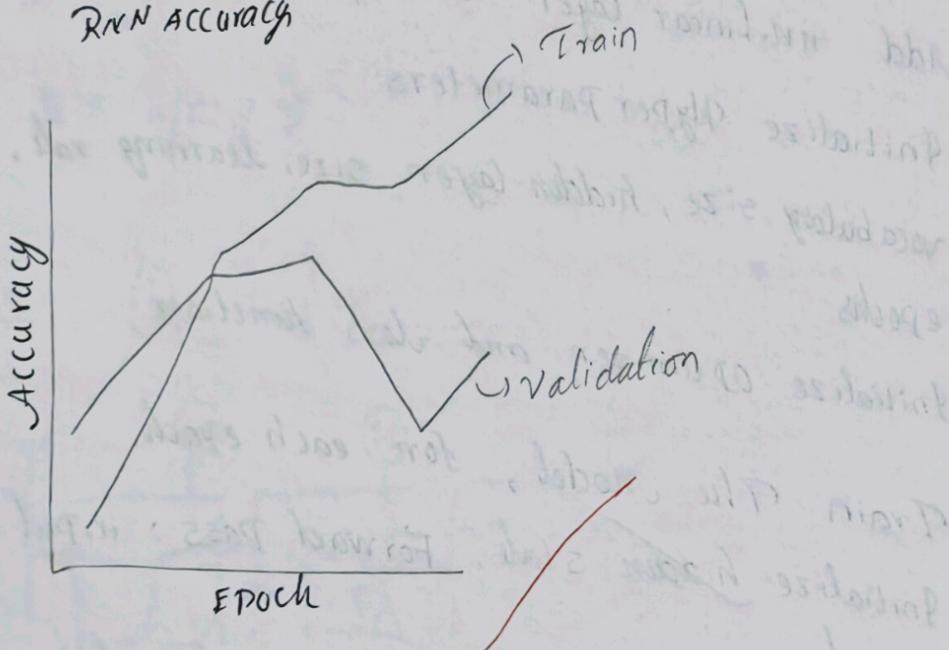
EPOCH 3
Accuracy: 0.8231 - Loss: 0.4031

EPOCH 4
Accuracy: 0.8235 - Loss: 0.4016

EPOCH 5
Accuracy: 0.8476 - Loss: 0.3546

RNN Test Accuracy = 68.15%, Loss: 0.5839

RNN Accuracy



due exp 09/25 8. Implement a LSTM

Aim: To implement a Long short Term memory model using Pytorch for time series forecasting.

Algorithm:

- * Data Preparation - Generate a sine wave dataset as the time series input.
- * Divide the sequence into input-output pairs
- * Model Design - Define an LSTM network
- * Use MSE as the loss function
- * Use Adam optimizer for parameter updates
- * Training Phase - initialize hidden and all states
- * compute loss between Predicted and Actual values
- * Backpropagate the error and update weights
- * Feed test input into the trained model
- * Plot Actual data vs Predicted data to evaluate Performance.

* Pseudocode:

BEGIN

Import libraries (torch, numpy, matplotlib)

Generate sine wave dataset

- Define timestamps

- Create input-output sequences with lookback window (seq_length)

Convert dataset to tensors

- X shape = (batch_size, seq_length, 1)

- Y shape = (batch_size, 1)

Define LSTMModel class

- LSTM layer (input_size, hidden_size, num_layers)

- Fully connected layer for output

- Forward Pass:

- Initialize h_0, c_0

- Pass input through LSTM

- Take last timestep output \rightarrow Pass through linear layer

- Initialize model, loss function (MSE), optimizer

- TRAINING LOOP for each epoch:

- Forward Pass: Predict $y_{\text{Pred}} = \text{model}(x)$

- Compute loss = MSE (y_{Pred}, y)

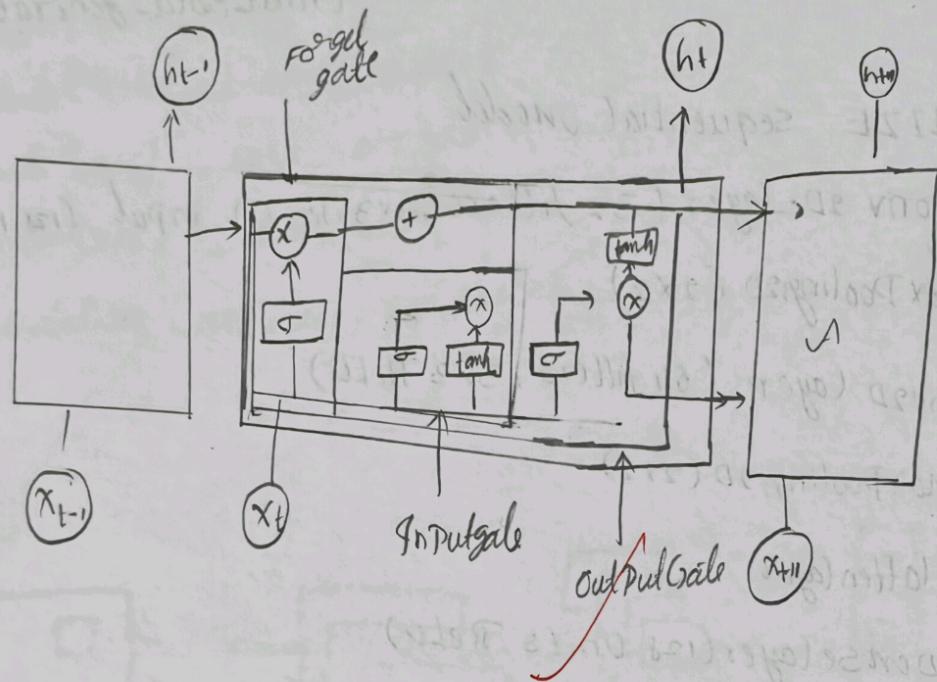
- Backpropagation: loss.backward()

Plot actual data vs Predicted data

Obs? Result?

END

LSTM Architecture



EPOCH 1

val accuracy : 0.6676 - val loss : 0.4151

EPOCH 2

val-accuracy : 0.8358 - val-loss : 0.3951

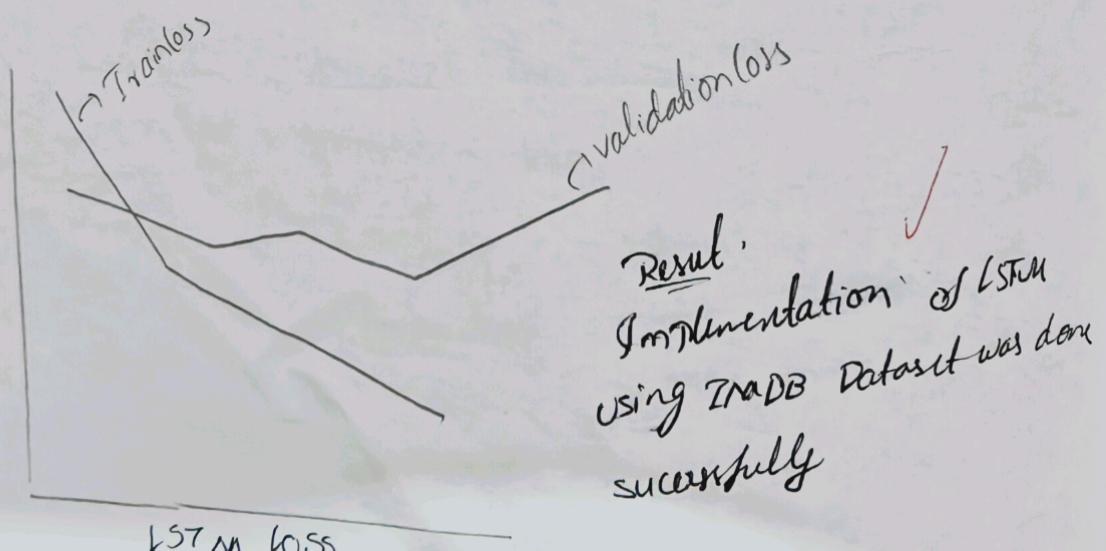
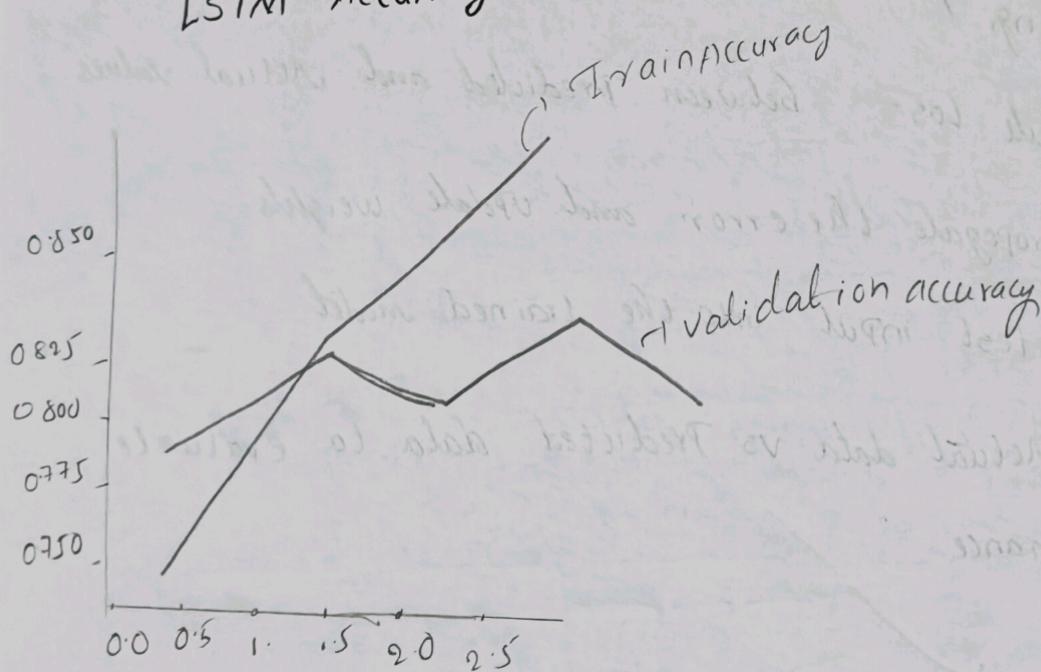
EPOCH 3

val-accuracy : 0.8328 - val-loss : 0.3862

EPOCH 4

val accuracy : 0.8528 - val loss : 0.3639

LSTM ACCURACY



```

# RNN Implementation
import tensorflow as tf
from tensorflow.keras.datasets import imdb
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, SimpleRNN, Dense, Dropout
import matplotlib.pyplot as plt

# Load IMDB dataset
vocab_size = 10000
 maxlen = 200

(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=vocab_size)
x_train = pad_sequences(x_train, maxlen=maxlen)
x_test = pad_sequences(x_test, maxlen=maxlen)

# Build RNN model
rnn_model = Sequential([
    Embedding(input_dim=vocab_size, output_dim=128, input_length=maxlen),
    SimpleRNN(128, dropout=0.2),
    Dense(1, activation='sigmoid')
])
rnn_model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

# Train model
rnn_history = rnn_model.fit(
    x_train, y_train,
    epochs=5,
    batch_size=64,
    validation_split=0.2
)

# Evaluate model
rnn_score = rnn_model.evaluate(x_test, y_test)
print(f"RNN Test Accuracy: {rnn_score[1]*100:.2f}% | Loss: {rnn_score[0]:.4f}")

# Plot Accuracy and Loss
plt.figure(figsize=(12,5))

# Accuracy
plt.subplot(1,2,1)
plt.plot(rnn_history.history['accuracy'], label='Train Accuracy')
plt.plot(rnn_history.history['val_accuracy'], label='Validation Accuracy')
plt.title('RNN Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

# Loss
plt.subplot(1,2,2)
plt.plot(rnn_history.history['loss'], label='Train Loss')
plt.plot(rnn_history.history['val_loss'], label='Validation Loss')
plt.title('RNN Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

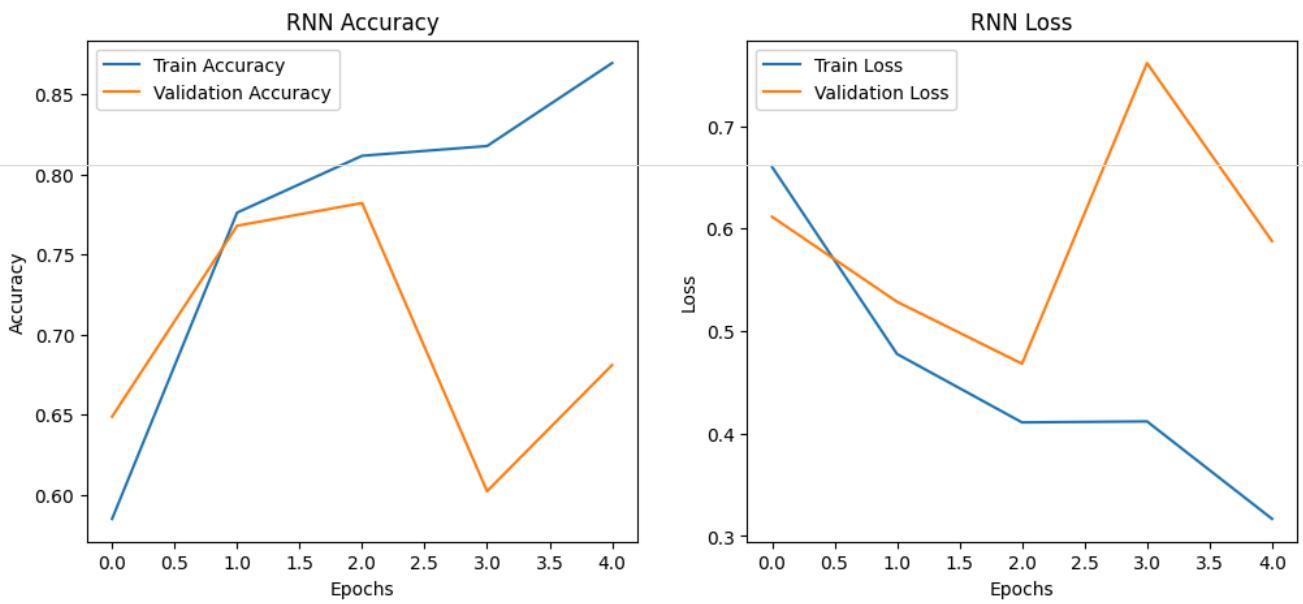
plt.show()

```

```

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb.npz
17464789/17464789 0s 0us/step
Epoch 1/5
/usr/local/lib/python3.12/dist-packages/keras/src/layers/core/embedding.py:97: UserWarning: Argument `input_length` is deprecated
  warnings.warn(
313/313 48s 147ms/step - accuracy: 0.5433 - loss: 0.6817 - val_accuracy: 0.6488 - val_loss: 0.6114
Epoch 2/5
313/313 80s 141ms/step - accuracy: 0.7572 - loss: 0.5022 - val_accuracy: 0.7680 - val_loss: 0.5285
Epoch 3/5
313/313 44s 140ms/step - accuracy: 0.8231 - loss: 0.4031 - val_accuracy: 0.7822 - val_loss: 0.4681
Epoch 4/5
313/313 45s 143ms/step - accuracy: 0.8235 - loss: 0.4016 - val_accuracy: 0.6022 - val_loss: 0.7613
Epoch 5/5
313/313 81s 140ms/step - accuracy: 0.8476 - loss: 0.3546 - val_accuracy: 0.6810 - val_loss: 0.5877
782/782 18s 23ms/step - accuracy: 0.6824 - loss: 0.5871
RNN Test Accuracy: 68.15% | Loss: 0.5839

```



```

# LSTM Implementation
import tensorflow as tf
from tensorflow.keras.datasets import imdb
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, LSTM, Dense, Dropout
import matplotlib.pyplot as plt

# Load IMDB dataset
vocab_size = 10000
maxlen = 200

(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=vocab_size)
x_train = pad_sequences(x_train, maxlen=maxlen)
x_test = pad_sequences(x_test, maxlen=maxlen)

# Build LSTM model
lstm_model = Sequential([
    Embedding(input_dim=vocab_size, output_dim=128, input_length=maxlen),
    LSTM(128, dropout=0.2, recurrent_dropout=0.2),
    Dense(1, activation='sigmoid')
])
lstm_model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

# Train model
lstm_history = lstm_model.fit(
    x_train, y_train,
    epochs=5,
    batch_size=64,
    validation_split=0.2
)

# Evaluate model
lstm_score = lstm_model.evaluate(x_test, y_test)
print(f'LSTM Test Accuracy: {lstm_score[1]*100:.2f}% | Loss: {lstm_score[0]:.4f}')

# Plot Accuracy and Loss
plt.figure(figsize=(12,5))

# Accuracy
plt.subplot(1,2,1)
plt.plot(lstm_history.history['accuracy'], label='Train Accuracy')
plt.plot(lstm_history.history['val_accuracy'], label='Validation Accuracy')
plt.title('LSTM Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

# Loss
plt.subplot(1,2,2)
plt.plot(lstm_history.history['loss'], label='Train Loss')
plt.plot(lstm_history.history['val_loss'], label='Validation Loss')
plt.title('LSTM Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()

```

```

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb.npz
17464789/17464789 0s 0us/step
Epoch 1/5
/usr/local/lib/python3.12/dist-packages/keras/src/layers/core/embedding.py:97: UserWarning: Argument `input_length` is deprecated. Use `input_shape` instead.
  warnings.warn(
313/313 193s 603ms/step - accuracy: 0.6676 - loss: 0.5828 - val_accuracy: 0.8122 - val_loss: 0.4013
Epoch 2/5
313/313 198s 593ms/step - accuracy: 0.8517 - loss: 0.3529 - val_accuracy: 0.8358 - val_loss: 0.3946
Epoch 3/5
313/313 206s 606ms/step - accuracy: 0.8840 - loss: 0.2872 - val_accuracy: 0.8328 - val_loss: 0.3946
Epoch 4/5
313/313 204s 611ms/step - accuracy: 0.9019 - loss: 0.2525 - val_accuracy: 0.8528 - val_loss: 0.3946
Epoch 5/5
313/313 198s 598ms/step - accuracy: 0.9328 - loss: 0.1832 - val_accuracy: 0.8338 - val_loss: 0.3946
782/782 62s 79ms/step - accuracy: 0.8309 - loss: 0.4013
LSTM Test Accuracy: 83.48% | Loss: 0.3946

```

